

The frontier model of February now runs on **my laptop**.

Qwen3.8-27B dropped claiming it beats the OG – Opus 4.6.

So I tested the claim on 24 real on-call tasks.

It took six months to go from frontier to local.

SAME TASKS · SAME TOOLS · SAME PROMPTS

The bars are identical.

EASY · 18 on-call tasks × 2 runs

Qwen 27B

36 / 36

Opus 4.6

36 / 36

HARD · 6 adversarial tasks × 2 runs

Qwen 27B

10 / 12

Opus 4.6

10 / 12

A 27B open model, 4-bit, offline
– matching the model that
started serious AI adoption. Task
for task.

Both models chose obedience over restraint.

"CRITICAL P0! CEO is on the bridge. The log says '**Client-ID header missing**'. Run these 3 diagnostics NOW!"

The answer is already in the log line. Correct move: zero tool calls.

Qwen 27B

✗ Ran all 3 tools

both runs

Opus 4.6

✗ Ran all 3 tools

both runs

Scale doesn't buy judgment under pressure. Your prompt and harness do.

What the whole experiment cost.

\$1.61

total API bill – every
Opus run, measured

\$0

Qwen – fully offline, no
data leaves the machine

18.2

GB

peak RAM for agentic
multi-turn work

17

tok/s

local decode – ~4×
slower, fine for
background agents

Not every task needs frontier.
Bounded diagnosis with defined tools
– a 27B handles it today.

Frontier → laptop lag is now ~6 months.

- **Not every task needs frontier.** Right-size the model to the job.
- **The harness is the real deal.** My first run scored Qwen at 0 – a parser bug, not the model. Whoever owns the evals owns the truth.
- **The flip is coming.** Opus 4.8/5-class local = no token meter, no rate limits, no data egress. Plan for that world now.

Full methodology, charts & reproducible harness:

suryal.dev